SAMPLE DELIVERABLE / DISCOVERY & FEASIBILITY

Supplier quotation review workflow

An illustrative assessment showing how a document-heavy process can be scoped before any automation is built.

SCENARIO	ASSESSMENT TYPE	OUTPUT
Fictional mid-sized procurement team	Workflow and technical feasibility	Bounded pilot recommendation

Executive summary

The fictional team receives supplier quotations in PDF and spreadsheet formats. Staff manually compare line items, delivery terms, warranties and tax details before preparing a recommendation. The work is repetitive, but the decision still requires human judgement because quotations can be incomplete, ambiguous or commercially non-equivalent.

The recommended first step is not a fully autonomous purchasing agent. It is a bounded document-review pilot that extracts comparable fields, highlights inconsistencies, records evidence and leaves approval with a person.

RECOMMENDATION	**Pilot a review-assist workflow: extract → normalise → compare → flag exceptions → human approval.**

What this example is—and is not

This document is a fictional sample created to demonstrate assessment structure and engineering thinking. Company names, workflows, findings, timeline and examples are illustrative. It is not a client report, a statement of actual customer needs, a performance benchmark, a quote or a commitment to deliver within the example timeline.

01 / CURRENT-STATE REVIEW

Understand the work before choosing the technology.

Illustrative workflow observed in the scenario

STEP	CURRENT ACTIVITY	FRICTION / RISK
01	Receive quotation files by email and shared folders.	File names and formats vary; duplicate versions can be hard to spot.
02	Read supplier details, prices, units and terms.	Important values may be embedded in tables, footnotes or scanned pages.
03	Copy values into a comparison spreadsheet.	Manual rekeying creates transcription and unit-normalisation risk.
04	Check missing warranty, delivery and tax details.	Exceptions are not always consistently documented.
05	Prepare a recommendation for an authorised reviewer.	The rationale and source evidence may be separated from the summary.

Questions to resolve during a real discovery

• Which document formats and languages are in scope, and how often do scanned files appear?

• Which fields are mandatory for a fair comparison, and which terms require human interpretation?

• Are currencies, units, taxes and delivery terms normalised by an existing policy?

• Who can review exceptions, approve a recommendation and override an extracted value?

• Where may files be stored and processed? Are there retention, residency or vendor restrictions?

Assessment principle

Separate extraction from decision-making. A model may propose a value and cite the source location, but deterministic checks should handle known rules and a named reviewer should approve commercially significant decisions.

02 / TARGET DESIGN

A review-assist system with visible control points.

Proposed pilot architecture

LAYER	RESPONSIBILITY	CONTROL
Intake	Accept approved PDF/XLSX files and record document metadata.	File-type checks, size limits, access control and malware scanning where available.
Extraction	Extract supplier identity, line items, units, prices and terms.	Store source page/row references and confidence/review indicators.
Normalisation	Convert values to a comparison schema and consistent units.	Deterministic validation for currencies, units, tax and required fields.
Comparison	Compare like-for-like costs and flag non-equivalent terms.	Rules are explicit, versioned and independently testable.
Review	Present evidence, exceptions and a draft summary to an authorised user.	Human approval required; reviewer actions are recorded.

Reliability and governance controls

• **Evidence:** Each extracted field links to a page, table row or source span so a reviewer can verify it.

• **Exceptions:** Missing fields, conflicting values, low-confidence extraction and unsupported formats route to manual review.

• **Security:** Access is role-based; uploads and outputs follow an agreed retention policy; secrets stay server-side.

• **Evaluation:** A representative, permissioned test set measures field-level accuracy, exception recall and reviewer corrections.

• **Fallback:** The manual process remains available if extraction, model or integration services fail.

What is deliberately not automated

The pilot does not award contracts, send purchase orders, negotiate with suppliers or approve spend. Those actions remain outside the automated workflow unless separately assessed, authorised and controlled.

03 / PILOT PLAN & DECISION

Prove the narrow workflow before expanding it.

Illustrative delivery plan

STAGE	WORK PRODUCT	EXIT CHECK
Discovery	Process map, sample review, constraints and agreed field schema.	Representative files and review rules are approved.
Pilot build	Upload, extraction, comparison view, evidence links and review actions.	Core workflow runs on an agreed test set.
Evaluation	Error analysis, exception review, usability feedback and risk log.	Acceptance thresholds and known limitations are documented.
Handover decision	Deployment notes, operator guide and next-phase recommendation.	Client chooses to stop, refine or scope a production phase.

Example acceptance measures to agree—not results

- Field extraction quality measured per field and document type, not as one blended accuracy number.
- Critical exception recall for missing warranty, unit mismatch, tax mismatch and delivery-term conflicts.
- Percentage of generated comparisons with verifiable source references.
- Reviewer correction rate and time required to complete the review, compared with a documented baseline.

Risks to resolve before production

RISK	MITIGATION
Poor scans or unusual layouts	Define supported formats; route unsupported files to manual processing.
Incorrect equivalence assumptions	Approve the comparison schema and rules with procurement stakeholders.
Sensitive supplier data	Agree data handling, provider selection, retention and access before using real files.
Model/provider downtime	Retain manual fallback and make failed or partial runs visible.

Next step

A real assessment would begin with a short conversation, then confirm representative files, users, constraints and the decision that needs to improve. No build is started until scope, fee, data handling and acceptance criteria are agreed in writing.

This fictional report is a sample of the format and thinking—not a quote, a guarantee or a claim of prior client work.